

# EXTENDING TENSORRT WITH CUSTOM LAYERS

Dr. Pallab Maji, Senior Solutions Architect – Deep Learning



# CPP & CUDA SETUP

# Prepare the CUDA Kernel

Prepare and Verify the CUDA Kernel for the Layer to be Replaced

```
geluPluginv2 > geluPlugin.cu
25 // constants for approximating the normal cdf
26 constexpr float A = 0.5;
27 constexpr float B = 0.7978845608028654; // sqrt(2.0/M_PI)
28 constexpr float C = 0.035677408136300125; // 0.044715 * sqrt(2.0/M_PI)
29
30 template <typename T, unsigned TPB>
31 __global__ void geluKernel(const T a, const T b, const T c, int n, const T* input, T* output)
32 {
33     const int idx = blockIdx.x * TPB + threadIdx.x;
34
35     if (idx < n)
36     {
37         const T in = input[idx];
38         const T cdf = a + a * tanh(in * (c * in * in + b));
39         output[idx] = in * cdf;
40     }
41 }
42
43
44 int computeGelu(cudaStream_t stream, int n, const float* input, float* output)
45 {
46     constexpr int blockSize = 256;
47     const int gridSize = (n + blockSize - 1) / blockSize;
48     geluKernel<float, blockSize><<<gridSize, blockSize, 0, stream>>>(A, B, C, n, input, output);
49
50     // CHECK(cudaPeekAtLastError());
51     return 0;
52 }
53 }
```

# custom plugin – Setup

Look into the CPP code of the sample plugin

```
using namespace nvinfer1;  —————> Namespace for plugin

// GELU plugin specific constants
namespace {
    static const char* GELU_PLUGIN_VERSION{"1"};  —————> Add plugin version
    static const char* GELU_PLUGIN_NAME{"CustomGeluPlugin"};  —————> Add plugin name
}

// Static class fields initialization
PluginFieldCollection GeluPluginCreator::mFC{};  —————> Set and Remember the plugin creator name
std::vector<PluginField> GeluPluginCreator::mPluginAttributes;

REGISTER_TENSORRT_PLUGIN(GeluPluginCreator);  —————> Register plugin creator name in plugin factory
```

# custom plugin - Constructors

Look into the CPP code of the sample plugin (contd.)

```
89  GeluPlugin::GeluPlugin(const std::string name, const DataType type)
90      : mLayerName(name)
91      , mType(type)
92      , mHasBias(false)
93      , mLd(0)
94  {
95      mBias.values = nullptr;
96      mBias.count = 0;
97  }
98
99  GeluPlugin::GeluPlugin(const std::string name, const DataType type, const Weights B)
100      : mLayerName(name)
101      , mType(type)
102      , mHasBias(true)
103      , mBias(B)
104      , mLd(B.count)
105  {
106  }
```

→ Mention all values that is used by the plugin in the constructor

# custom plugin - Serialize

Look into the CPP code of the sample plugin (contd.)

```
size_t GeluPlugin::getSerializationSize() const
```

```
{
```

```
    const size_t wordSize = getElementSize(mType);
```

```
    const size_t biasSize = mHasBias ? mLd * wordSize : 0;
```

```
    return sizeof(mType) + sizeof(mHasBias) + sizeof(mLd) + sizeof(mInputVolume) + biasSize;
```

```
}
```

Based on the variables to be saved, compute the size

```
void GeluPlugin::serialize(void* buffer) const
```

```
{
```

```
    serialize_value(&buffer, mHasBias);
```

```
    serialize_value(&buffer, mInputVolume);
```

```
    serialize_value(&buffer, mType);
```

```
    serialize_value(&buffer, mLd);
```

Serialize will help in saving tensorrt optimized model

Mention all values that is used by the plugin constructor that needs to be saved

## custom plugin – Constructor to DE-Serialize

Look into the CPP code of the sample plugin (contd.)

```
GeluPlugin::GeluPlugin(const std::string name, const void* data, size_t length)
```

```
: mLayerName(name)
```



Constructor to load  
serialized plugin


```
{  
    std::cout << "Starting to deserialize GELU plugin" << std::endl;
```

```
    deserialize_value(&data, &length, &mHasBias);
```

```
    deserialize_value(&data, &length, &mInputVolume);
```

```
    deserialize_value(&data, &length, &mType);
```

```
    deserialize_value(&data, &length, &mLd);
```



Mention all values that is used by  
the plugin that was saved in serialize

# custom plugin – Constructor to DE-Serialize

Look into the CPP code of the sample plugin (contd.)

```
GeluPlugin::GeluPlugin(const std::string name, const void* data, size_t length)
```

```
: mLayerName(name)
```

```
{
```

```
std::cout << "Starting to deserialize GELU plugin" << std::endl;
```

```
deserialize_value(&data, &length, &mHasBias);
```

```
deserialize_value(&data, &length, &mInputVolume);
```

```
deserialize_value(&data, &length, &mType);
```

```
deserialize_value(&data, &length, &mLd);
```

Constructor to load  
serialized plugin

Mention all values that is used by  
the plugin that was saved in serialize



# custom plugin – ENQUEUE Function

Look into the CPP code of the sample plugin (contd.)

```
int GeluPlugin::enqueue(int batchSize, const void* const* inputs, void** outputs, void*, cudaStream_t stream)
{
    const int inputVolume = mInputVolume;
    int status = -1;

    // This plugin outputs only one tensor
    // Launch CUDA kernel wrapper and save its return value
    if (mType == DataType::kFLOAT)
    {
        const float* input = static_cast<const float*>(inputs[0]);
        float* output = static_cast<float*>(outputs[0]);
        if (mHasBias)
        {
            const float* bias = reinterpret_cast<float*>(mBiasDev);
            const int cols = inputVolume / mLd;
            const int rows = mLd;
            computeGeluBias(output, input, bias, rows, cols, stream);
        }
        else
        {
            status = computeGelu(stream, inputVolume, input, output);
        }
    }
    else
    {
        assert(false);
    }

    return status;
}
```

This function runs during the  
TensorRT Runtime.

Launch the verified CUDA  
kernel from here based  
on the condition

# BUILD Tensorrt Plugin with CMAKE

Build the plugin to generate a library then verify in Python

```
import ctypes

# Add plugin compiled library
ctypes.CDLL("../build/libGeluPlugin.so")
```



# PYTHON SETUP

# Freeze the Graph

Load saved model, Remove Training Nodes, Convert Variables to Constants & Save to Disk

```
# First freeze the graph and remove training nodes.
output_names = model.output.op.name # output_name is "dense_2/MatMul" for verification
sess = get_session()
frozen_graph = tf.graph_util.convert_variables_to_constants(
    sess, sess.graph.as_graph_def(), [output_names])
frozen_graph = tf.graph_util.remove_training_nodes(frozen_graph)
# Save the model
with open(frozen_filename, "wb") as fptr:
    fptr.write(frozen_graph.SerializeToString())
```

→ Convert Variables to Constants

→ Remove training nodes

# Convert FROZEN File to UFF

```
# Add plugin compiled library
```

```
ctypes.CDLL("../build/libGeluPlugin.so")
```

Load Plugin Compiled Library

```
def create_plugin_node(dynamic_graph):
```

```
    gelu_node = gs.create_plugin_node(  
        name="GeluActivation", op="CustomGeluPlugin", typeId=0)
```

Add a node with OP Name same as in the  
CPP file

```
    namespace_plugin_map = {"GeluActivation": gelu_node}
```

```
    dynamic_graph.collapse_namespaces(namespace_plugin_map)
```

Replace the node in the graph  
with custom TRT plugin

```
# Transform graph using graphsurgeon to map unsupported TensorFlow
```

```
# operations to appropriate TensorRT custom layer plugins
```

```
dynamic_graph = gs.DynamicGraph(frozen_graph)
```

```
create_plugin_node(dynamic_graph)
```

Load graph as dynamic graph using  
graphsurgeon

```
uff_model = uff.from_tensorflow(dynamic_graph, [output_names])
```

Convert to UFF

```
with open(uff_filename, "wb") as fptr:
```

```
    fptr.write(uff_model)
```

Save to Disk

# Build & Infer

Build the TensorRT engine as usual and do Inference

```
def build_engine(model_file, TRT_LOGGER):  
    # For more information on TRT basics, refer to the introductory samples.  
    with trt.Builder(TRT_LOGGER) as builder, builder.create_network() \  
        as network, trt.UffParser() as parser:  
        builder.max_workspace_size = 1 << 16  
        # Parse the Uff Network  
        parser.register_input("input_1", (3, 150, 150))  
        parser.register_output('dense_2/Softmax')  
        parser.parse(model_file, network)  
        # Build and return an engine.  
        return builder.build_cuda_engine(network)
```